

Lecture 20 R scripts

Example: Variational Bayes applied to farm data.

To illustrate Variational Bayes, we consider an agricultural dataset, which can be downloaded as `farmdata.csv` (for weaning weights) and `Kmat.csv` (for the kinship matrix) from Canvas. To benchmark Variational Bayes, we will compare results to those obtained using Gibbs sampling.

Writing a Gibbs sampler for this problem The Gibbs sampler for this problem is just a modification of the Gibbs sampler written in lab 7. Now however we can work with arbitrary variance-covariance matrices in the normal prior by including the argument `Kinv`. $K_{\text{inv}} = \mathbf{K}^{-1}$ is the inverse of the prior variance-covariance matrix up to proportionality.

```
#Arguments are
#iter: no of iterations
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#burnin: number of initial iterations to discard.
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#Kinv: inverse of K, where p(u) = N(0, \sigma^2_u K)
#a.u, b.u: hyper-parameters of gamma prior for tauu
#a.e, b.e: hyper-parameters of gamma prior for taue

normalmm.Gibbs<-function(iter,Z,X,y,burnin,taue_0,tauu_0,Kinv,a.u,b.u,a.e,b.e){
  n  <-length(y) #no. observations
  p  <-dim(X)[2] #no of fixed effect predictors.
  q  <-dim(Z)[2] #no of random effect levels
  tauu<-tauu_0
  taue<-taue_0
  beta0<-rnorm(p)
  u0   <-rnorm(q,0,sd=1/sqrt(tauu))

  #Building combined predictor matrix.
  W<-cbind(X,Z)
  WTW <-crossprod(W)
  WTy <-crossprod(W,y)
  library(mvtnorm)

  #storing results.
  par <-matrix(0,iter,p+q+2)
  #Calculating log predictive densities
  lppd<-matrix(0,iter,n)

  #Create modified identity matrix for joint posterior.
```

```

I0 <-diag(p+q)
diag(I0)[1:p]<-0
I0[-c(1:p),-c(1:p)] <-Kinv

for(i in 1:iter){
  #Conditional posteriors.
  uKinvu <- t(u0)%*%Kinv%*%u0
  uKinvu <-as.numeric(uKinvu)
  tauu <-rgamma(1,a.u+0.5*q,b.u+0.5*uKinvu)
  #Updating component of normal posterior for beta,u
  Prec <-WTW + tauu*I0/taue
  P.mean<- solve(Prec)%*%WTy
  P.var <-solve(Prec)/taue
  betau <-rmvnorm(1,mean=P.mean,sigma=P.var)
  betau <-as.numeric(betau)
  err <- y-W%*%betau
  taue <-rgamma(1,a.e+0.5*n,b.e+0.5*sum(err^2))
  #storing iterations.
  par[i,]<-c(betau,1/sqrt(tauu),1/sqrt(taue))
  beta0 <-betau[1:p]
  u0 <-betau[p+1:q]
  lppd[i,]= dnorm(y,mean=as.numeric(W%*%betau),sd=1/sqrt(taue))
}

lppd      = lppd[-c(1:burnin),]
lppdest   = sum(log(colMeans(lppd)))      #Estimating lppd for whole dataset.
pwaic2    = sum(apply(log(lppd),2,FUN=var)) #Estimating effective number of parameters.
par <-par[-c(1:burnin),]
colnames(par)<-c(paste('beta',1:p,sep=''),paste('u',1:q,sep=''),'sigma_b','sigma_e')
mresult<-list(par,lppdest,pwaic2)
names(mresult)<-c('par','lppd','pwaic')
return(mresult)
}

```

```

data<-read.csv('./farmdata.csv')
# data<-read.csv(file.choose()) #Choose farmdata.csv as downloaded from Canvas
K <-read.csv('./Kmat.csv')
# K<-read.csv(file.choose()) #Choose farmdata.csv as downloaded from Canvas
K <-as.matrix(K)
Kinv<-solve(K)

n<-dim(data)[1]
q<-dim(Kinv)[1]
X<-table(1:n,data$flock) #flock is fixed effect
#Indicator matrix for parents.
Z2<-table(1:n,data$sire)
Z3<-cbind(Z2,table(1:n,data$dam))

```

Importing and formatting the data

```
system.time(chain1<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=5,tauu_0=0.2,
                                   Kinv=Kinv,a.u=0.001,b.u=0.001,a.e=0.001,b.e=0.001))
```

running the Gibbs sampler, checking convergence, and calculating effective sample size

```
## Warning: package 'mvtnorm' was built under R version 4.3.1
```

```
##      user  system elapsed
##    1.741    0.029    1.772
```

```
system.time(chain2<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=1,tauu_0=1,
                                   Kinv=Kinv,a.u=0.001,b.u=0.001,a.e=0.001,b.e=0.001))
```

```
##      user  system elapsed
##    1.683    0.017    1.700
```

```
system.time(chain3<-normalmm.Gibbs(iter=10000,Z=Z3,X=X,y=data$y,burnin=2000,taue_0=0.2,tauu_0=3,
                                   Kinv=Kinv,a.u=0.001,b.u=0.001,a.e=0.001,b.e=0.001))
```

```
##      user  system elapsed
##    1.692    0.019    1.712
```

```
library(coda)
```

```
## Warning: package 'coda' was built under R version 4.3.1
```

```
rml1<-as.mcmc.list(as.mcmc((chain1$par[1:4000,])))
rml2<-as.mcmc.list(as.mcmc((chain2$par[1:4000,])))
rml3<-as.mcmc.list(as.mcmc((chain3$par[1:4000,])))
rml4<-as.mcmc.list(as.mcmc((chain1$par[4000+1:4000,])))
rml5<-as.mcmc.list(as.mcmc((chain2$par[4000+1:4000,])))
rml6<-as.mcmc.list(as.mcmc((chain3$par[4000+1:4000,])))
rml<-c(rml1,rml2,rml3,rml4,rml5,rml6)
```

```
#Gelman-Rubin diagnostic.
gelman.diag(rml)[[1]]
```

```
##      Point est. Upper C.I.
## beta1    1.0001886  1.000618
## beta2    1.0003186  1.000664
## u1       1.0002849  1.000427
## u2       1.0004086  1.000672
## u3       1.0002609  1.000783
## u4       1.0002456  1.000754
## u5       1.0001370  1.000347
## u6       1.0001749  1.000684
## u7       1.0001805  1.000262
## u8       1.0000731  1.000364
```

```
## u9      1.0001241  1.000346
## u10     1.0000389  1.000371
## u11     0.9999831  1.000121
## u12     1.0000541  1.000331
## u13     1.0000604  1.000143
## sigma_b 1.0009851  1.002606
## sigma_e 1.0012429  1.002924
```

```
#effective sample size.
effectiveSize(rml)
```

```
##      beta1      beta2      u1      u2      u3      u4      u5      u6
## 24000.000 24000.000 24000.000 21692.020 23909.883 24000.000 23848.486 24000.000
##      u7      u8      u9      u10     u11     u12     u13     sigma_b
## 24000.000 25176.790 23691.012 23376.080 24000.000 24000.000 24000.000 10917.902
##      sigma_e
##      8676.271
```

Writing CAVI algorithm to find V-B estimates To determine the parameters of the Variational Bayes approximation, we will use an CAVI algorithm. This means we will succesively update $E(\beta)$, $\text{Var}(\beta)$, $E(\mathbf{u})$, $\text{Var}(\mathbf{u})$, $E(\tau_u)$, $E(\tau_e)$ until convergence has been reached. In this code, convergence is defined as when $\sqrt{(\theta^{(t)} - \theta^{(t-1)})^2 / (|\theta^{(t)}| + 0.01)} \leq \epsilon$ for all elements of θ . The addition of 0.01 is to avoid problems when $\theta^{(t)} = 0$

```
#Arguments are
#epsilon: accuracy cut-off.
#iter: no of iterations
#Kinv: inverse of K, where p(u) = N(0, \sigma^2_u K)
#Z: Predictor matrix for random effects
#X: Predictor matrix for fixed effects
#y: response vector
#taue_0: initial guess for residual precision.
#tauu_0: initial guess for random effect precision
#a.u, g.u: hyper-parameters of gamma prior for tauu
#a.e, g.e: hyper-parameters of gamma prior for taue

#Output are final estimates, plus iteration number when convergence was reached.
VB.mm<-function(epsilon,iter,Kinv,Z,X,y,taue_0,tauu_0,u0,beta0,a.e,g.e,a.u,g.u){
  n<-dim(X)[1]
  p<-dim(X)[2]
  q<-dim(Z)[2]
  ZTZ<-crossprod(Z)
  ZTY<-crossprod(Z,y)
  XTY<-crossprod(X,y)
  ZTX<-crossprod(Z,X)
  XTX<-crossprod(X)
  XTXinv<-solve(XTX)

  for(i in 1:iter){
    Vu <-solve(taue_0*ZTZ+tauu_0*Kinv) #update Var(u)
    u  <-taue_0*Vu%*%(ZTY-ZTX%*%beta0) #update E(u)
    Vb <-XTXinv/taue_0 #update Var(beta)
    b  <-XTXinv%*%(XTY-t(ZTX)%*%u) #update E(beta)
```

```

TrKinvu <- sum(diag(Kinv%*%Vu))
uKinvu <- t(u)%*%Kinv%*%u
tauu <- (a.u+0.5*q)/(g.u+0.5*as.numeric(uKinvu)+0.5*TrKinvu)
tauu <- as.numeric(tauu)
err <- y - X%*%b - Z%*%u
TrXTXb <- sum(diag(CTX%*%Vb))
TrZTZu <- sum(diag(ZTZ%*%Vu))
taue <- (a.e+0.5*n)/(g.e+0.5*sum(err^2)+0.5*TrXTXb+0.5*TrZTZu)
taue <- as.numeric(taue)

if(i > 1){
diffb <- sqrt((b-beta0)^2)/(abs(b)+0.01)
diffu <- sqrt((u-u0)^2)/(abs(u)+0.01)
diffte <- abs(taue_0-taue)/(taue+0.01)
difftu <- abs(tauu_0-tauu)/(tauu+0.01)
diffvb <- sqrt((diag(Vb0) - diag(Vb))^2)/(diag(Vb))
diffvu <- sqrt((diag(Vu0) - diag(Vu))^2)/(diag(Vu))
diff.all<-c(diffb,diffu,diffte,difftu,diffvb,diffvu)
if(max(diff.all) < epsilon) break
}
Vu0 <- Vu;u0<-u;Vb0<-Vb;beta0<-b;taue_0<-taue;tauu_0<-tauu
#Calculate relative change.
}

taue.param<-c((a.e+0.5*n),(g.e+0.5*sum(err^2)+0.5*TrXTXb+0.5*TrZTZu))
tauu.param<-c((a.u+0.5*q),(g.u+0.5*uKinvu+0.5*TrKinvu))
param<-list(b,Vb,u,Vu,taue.param,tauu.param,i)
names(param)<-c('beta_mean','beta_var','u_mean','u_var','tau_e','tau_u','iter')
return(param)
}

system.time(test1<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=0.2,tauu_0=0.2,
u0=rnorm(13),beta0=rnorm(2),a.e=0.001,g.e=0.001,a.u=0.001,g.u=0.001))

```

Estimating parameter of variational Bayes approximations

```
## user system elapsed
## 0.093 0.000 0.094
```

```
system.time(test2<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=5,tauu_0=1,
u0=rnorm(13),beta0=rnorm(2),a.e=0.001,g.e=0.001,a.u=0.001,g.u=0.001))
```

```
## user system elapsed
## 0.074 0.001 0.075
```

```
system.time(test3<-VB.mm(epsilon=1e-5,iter=2000,Kinv=Kinv,Z=Z3,X=X,y=data$y,taue_0=1,tauu_0=5,
u0=rnorm(13),beta0=rnorm(2),a.e=0.001,g.e=0.001,a.u=0.001,g.u=0.001))
```

```
## user system elapsed
## 0.055 0.001 0.056
```

```
test1$iter
```

```
## [1] 1070
```

```
test2$iter
```

```
## [1] 1152
```

```
test3$iter
```

```
## [1] 870
```

```
#Comparing point estimates/posterior means
```

```
chain.all<-rbind(chain1$par,chain2$par,chain3$par)
```

```
#beta
```

```
cbind(test1$beta_mean,test2$beta_mean,test3$beta_mean,colMeans(chain.all[,1:2]))
```

Comparing estimates obtained from Gibbs sampling and Variational Bayes

```
##      [,1]      [,2]      [,3]      [,4]
```

```
## 1 4.149860 4.149861 4.149861 4.153120
```

```
## 2 4.708765 4.708766 4.708767 4.709322
```

```
#u
```

```
cbind(test1$u_mean,test2$u_mean,test3$u_mean,colMeans(chain.all[,3:15]))
```

```
##      [,1]      [,2]      [,3]      [,4]
```

```
## 1  0.80578542  0.80578469  0.80578421  0.80841251
```

```
## 2 -1.08621855 -1.08621928 -1.08621975 -1.09169227
```

```
## 3  0.21245313  0.21245245  0.21245200  0.21098242
```

```
## 4  0.47527965  0.47527943  0.47527928  0.47601605
```

```
## 5 -0.52479170 -0.52479204 -0.52479226 -0.53454094
```

```
## 6 -0.69870349 -0.69870374 -0.69870390 -0.70499465
```

```
## 7  0.50963975  0.50963940  0.50963918  0.51611965
```

```
## 8 -0.16302215 -0.16302238 -0.16302253 -0.16250230
```

```
## 9 -0.15022031 -0.15022065 -0.15022087 -0.14939405
```

```
## 10 0.12955976 0.12955948 0.12955930 0.13095244
```

```
## 11 0.14797212 0.14797175 0.14797150 0.15130713
```

```
## 12 -0.07160937 -0.07160964 -0.07160982 -0.07495579
```

```
## 13 0.57213111 0.57213069 0.57213042 0.57072883
```

```
#means of Variance parameters: E(\sigma_e^2), E(\sigma_u^2)
```

```
sigmae2<-sigmau2<-rep(0,3)
```

```
sigmae2[1]<-test1$tau_e[2]/(test1$tau_e[1]-1)
```

```
sigmau2[1]<-test1$tau_u[2]/(test1$tau_u[1]-1)
```

```
sigmae2[2]<-test2$tau_e[2]/(test2$tau_e[1]-1)
```

```
sigmau2[2]<-test2$tau_u[2]/(test2$tau_u[1]-1)
```

```
sigmae2[3]<-test3$tau_e[2]/(test3$tau_e[1]-1)
```

```
sigmau2[3]<-test3$tau_u[2]/(test3$tau_u[1]-1)
```

```
cbind(sigmau2,mean(chain.all[,16]^2))
```

```
##          sigmau2
## [1,] 0.352975 0.4649196
## [2,] 0.352975 0.4649196
## [3,] 0.352975 0.4649196
```

```
cbind(sigmae2,mean(chain.all[,17]^2))
```

```
##          sigmae2
## [1,] 0.05642792 0.05242605
## [2,] 0.05642792 0.05242605
## [3,] 0.05642792 0.05242605
```

```
#Comparing posterior distributions.
par(mfrow=c(5,4))
mlim<-quantile(chain.all[,1],c(0.005,0.995))
curve(dnorm(x,mean=test1$beta_mean[1],sd=sqrt(test1$beta_var[1,1])),ylab='Density',main='',
      xlim=mlim,col=2,lty=1,xlab=expression(beta[1]),cex.lab=2.5,cex.axis=1.5,lwd=2)
lines(density(chain.all[,1]),lwd=2)
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5,lwd=2)

mlim<-quantile(chain.all[,2],c(0.005,0.995))
curve(dnorm(x,mean=test1$beta_mean[2],sd=sqrt(test1$beta_var[2,2])),ylab='Density',main='',
      xlim=mlim,col=2,lty=1,xlab=expression(beta[2]),cex.lab=2.5,cex.axis=1.5,lwd=2)
lines(density(chain.all[,2]),lwd=2)
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5,lwd=2)

#Repeat for random effects.
for(i in 1:13){
  mlim<-quantile(chain.all[,i+2],c(0.005,0.995))
  curve(dnorm(x,mean=test1$u_mean[i],sd=sqrt(test1$u_var[i,i])),ylab='Density',main='',
        xlim=mlim,col=2,lty=1,xlab=paste('u',i,sep=''),cex.lab=2.5,cex.axis=1.5,lwd=2)
  lines(density(chain.all[,i+2]),lwd=2)
  legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5,lwd=2)
}

mlim<-quantile(chain.all[,16]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_u[1],rate=test1$tau_u[2])*x^(-2),ylab='Density',main='',
      xlim=c(0,mlim[2]),col=2,lty=1,xlab=expression(sigma[u]^2),cex.lab=2.5,cex.axis=1.5,lwd=2)
lines(density(chain.all[,16]^2),lwd=2)
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5,lwd=2)

mlim<-quantile(chain.all[,17]^2,c(0.005,0.995))
curve(dgamma(1/x,shape=test1$tau_e[1],rate=test1$tau_e[2])*x^(-2),ylab='Density',main='',lwd=2,
      xlim=c(0,mlim[2]),col=2,lty=1,xlab=expression(sigma[e]^2),cex.lab=2.5,cex.axis=1.5,cex.axis=1.5)
lines(density(chain.all[,17]^2),lwd=2)
legend('topright',legend=c('Gibbs','V-B'),col=1:2,lty=1,bty='n',cex=2.5,lwd=2)
```

