

Week 7 Lab Solutions – MAST90125: Bayesian Statistical learning

Question One

Consider a Poisson regression,

$$y_i \sim \text{Pois}(\lambda_i) \quad \text{and} \quad \log(\lambda_i) = \mathbf{x}_i' \boldsymbol{\beta}, \quad \boldsymbol{\beta} \in \mathbb{R}^p$$

In lectures we learned various techniques for approximating the posterior distribution. In this lab, attempt as many of these techniques as possible to complete the following tasks.

Consider the dataset `Warpbreaks.csv`, which can be downloaded from Canvas. This dataset contains information of the number of breaks in a consignment of wool. In addition, Wool type (A or B) and tension level (L, M or H) are recorded. To investigate the association between the number of breaks and wool type, various forms of generalised linear model are proposed where Bayesian computing techniques should be used.

As a reminder the following techniques will be considered for approximating the posterior distribution.

- Metropolis-Hastings algorithm.
- Gibbs sampler.

When coding, assume the prior for the coefficients $\boldsymbol{\beta} \sim N(\mathbf{0}, 5\mathbf{I}_p)$.

Some hints:

An initial guess can be determined from fitting a Poisson regression using the function `glm`. Treat wool type as a factor using the function `glm`

```
set.seed(123456)
warpbreak= read.csv(file = './warpbreaks.csv',header=TRUE)
#This line will need to be changed when you run this yourself.
mod<-glm(breaks~as.factor(wool),data=warpbreak,family='poisson')
summary(mod)
```

```
##
## Call:
## glm(formula = breaks ~ as.factor(wool), family = "poisson", data = warpbreak)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    3.43518    0.03454  99.443 < 2e-16 ***
## as.factor(wool)B -0.20599    0.05157  -3.994 6.49e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
```

```
##      Null deviance: 297.37  on 53  degrees of freedom
## Residual deviance: 281.33  on 52  degrees of freedom
## AIC: 560
##
## Number of Fisher Scoring iterations: 4
```

```
Sigma <-vcov(mod); Sigma
```

```
##              (Intercept) as.factor(wool)B
## (Intercept)    0.001193293    -0.001193293
## as.factor(wool)B -0.001193293     0.002659566
```

```
X<-model.matrix(mod)
```

Metropolis-Hastings code

```
#Part one: function for performing Metropolis-Hastings sampling for
#Poisson regression. Proposed transition distribution is normal,
#with mean = betahat and variance-covariance matrix c^2*Sigma.
#Namely,  $Q(\theta(t)|\theta(t-1)) = N(\theta(t)|\theta_{\text{tamean}}=\text{betahat}, c^2*Sigma)$ 
#Inputs:
#y: vector of responses
#X: predictor matrix including intercept.
#c: scalar associated with the variance-covariance matrix,
#thetamean: mean vector for the proposed transition distribution Q.
#Sigma: variance covariance matrix parameter in Q
#iter: number of iterations
#burnin: number of initial iterations to throw out.
library(mvtnorm)

## Warning: package 'mvtnorm' was built under R version 4.3.1

MetropolisHastings.fn<-function(y,X,c,thetamean,Sigma,iter,burnin){
  p <-dim(X)[2]      #number of parameters

  theta0<-rnorm(p)   #initial values
  thetas<-matrix(0,iter,p) #matrix to store values.
  thetas[1,]<-theta0
  indi<-0
  for(i in 1:(iter-1)){
    theta_t <-rmvnorm(1,mean=thetamean,sigma=c^2*Sigma) #draw candidate
    theta_t <-as.numeric(theta_t)
    xbc      <-X%*%theta_t
    p.c      <-exp(xbc) #Calculate lambda for candidate.
    xb       <-X%*%thetas[i,]
    p.b      <-exp(xb)   #Calculate lambda for theta(t-1)
    #Note for code to work correctly, sum goes over only the dpois part because dpois
    #evaluated over multiple observations is a vector but the dmnorm
    #for candidate/previous iteration is a scalar.
    r_up<-sum(dpois(y,lambda=p.c,log=TRUE))+sum(dnorm(theta_t,0,sqrt(5),log=TRUE)) +
      dmnorm(thetas[i,],mean=thetamean,sigma=c^2*Sigma,log=TRUE)
    #log joint dist + log proposal at the previous state
    r_low<-sum(dpois(y,lambda=p.b,log=TRUE))+sum(dnorm(thetas[i,],0,sqrt(5),log=TRUE)) +
      dmnorm(theta_t,mean=thetamean,sigma=c^2*Sigma,log=TRUE)
    #log joint dist + log proposal for the candidate state.
    r <-r_up-r_low #difference of log acceptance rate
    #Draw an indicator whether to accept/reject candidate
    ind<-rbinom(1,1,exp( min(c(r,0)) ) )
    thetas[i+1,]<- ind*theta_t + (1-ind)*thetas[i,]
    indi<-indi+ind
  }
  indi
  #discard initial iterations
  results<-thetas[-c(1:burnin),]
  names(results)<-c('beta0','beta1') #column names
  return(results)
}
```

```
#formatting data into the correct format.
```

```
#one way to determine betaest and Sigma
```

```
# betaest<-mod$coef
```

```
# Sigma=vcov(mod)
```

```
# mod$coef
```

```
# Sigma
```

```
#another way to determine betaest and Sigma
```

```
modest <-lm(log(breaks)~as.factor(wool),data=warpbreak)
```

```
betaest<-modest$coef
```

```
Sigma <-vcov(modest)
```

```
betaest
```

```
##      (Intercept) as.factor(wool)B
```

```
##      3.3174392      -0.1521536
```

```
Sigma
```

```
##              (Intercept) as.factor(wool)B
```

```
## (Intercept)      0.006982306      -0.006982306
```

```
## as.factor(wool)B -0.006982306      0.013964613
```

```
attempt2<-MetropolisHastings.fn(y=warpbreak$breaks,X=X,c=1,thetamean=betaest,Sigma=Sigma,  
                                iter=10000,burnin=2000)
```

```
#Posterior means
```

```
colMeans(attempt2)
```

```
## [1]  3.4298672 -0.2001627
```

```
#Posterior standard deviations
```

```
apply(attempt2,2,FUN=sd)
```

```
## [1] 0.03486872 0.05430271
```

```
#95 % central credible intervals
```

```
apply(attempt2,2,FUN=function(x) quantile(x,c(0.025,0.975)) )
```

```
##      [,1]      [,2]
```

```
## 2.5%  3.363289 -0.3020488
```

```
## 97.5% 3.501333 -0.1008965
```


```
par(mfrow=c(1,2))
```

```
#Plot marginal posteriors
```


```
plot(density(attempt2[,1]),type='l',xlab=expression(beta[0]),ylab='posterior density',  
      main='Metropolis-Hastings Algorithm')
```

```
plot(density(attempt2[,2]),type='l',xlab=expression(beta[1]),ylab='posterior density',  
      main='Metropolis-Hastings Algorithm')
```

Metropolis–Hastings Algorithm



Metropolis–Hastings Algorithm




```
length(unique(attempt2[,1]))
```

```
## [1] 699
```


```
#Plot the simulated beta samples
```

```
plot(attempt2[,1], xlab='Iteration', ylab=expression(beta[0]), type='l',  
     main=expression(paste(beta[0], ' sequence generated by MH'))))  
plot(attempt2[,2], xlab='Iteration', ylab=expression(beta[1]), type='l',  
     main=expression(paste(beta[1], ' sequence generated by MH'))))
```


β_0 sequence generated by MH



β_1 sequence generated by MH



```
#ACF plot
acf(attempt2[,1], main=expression(beta[0]), cex.main=2)
acf(attempt2[,2], main=expression(beta[1]), cex.main=2)
```



Gibbs sampler

It can be shown that the posterior pdf of $\beta = (\beta_0, \beta_1)'$ is

$$p(\beta_0, \beta_1 | (y_1, x_1), \dots, (y_n, x_n)) \propto \exp\left\{\left(\sum_i y_i\right)\beta_0 + \left(\sum_i y_i x_i\right)\beta_1 - 0.1(\beta_0^2 + \beta_1^2) - e_0^\beta \left(\sum_i e^{x_i \beta_1}\right)\right\}.$$

Thus, it is difficult to find the conditional posterior pdf of β_0 given β_1 and that of β_1 given β_0 . Therefore, Gibbs sampler is not a good method for the problem considered here.